

---

# Exploring Deep Learning Techniques for Long-Tailed Recognition: Methods, Models, and Analysis

MASTER'S THESIS IN  
ELECTRICAL ENGINEERING

---

By

**Christine Annelise Midtgaard**

AU521655

STUDY PROGRAM: ELECTRICAL ENGINEERING

SUPERVISOR

**Kim Bjerger**

ASSOCIATE PROFESSOR

kbe@ece.au.dk



M.Sc. IN ELECTRICAL ENGINEERING  
DEPARTMENT OF ENGINEERING SCIENCE & TECHNOLOGY  
AARHUS UNIVERSITY

JANUARY 3, 2025



# Abstract

This thesis, titled *Exploring Deep Learning Techniques for Long-Tailed Recognition: Methods, Models, and Analysis*, by Christine Annelise Midtgaard, explores the challenges and solutions related to training deep learning models on long-tailed datasets. Long-tailed datasets, where a few classes dominate with abundant samples while many classes have sparse representation, pose significant challenges for traditional training methods. These imbalances often lead to models that perform well on majority classes but struggle to recognize or generalize to minority (tail) classes.

This thesis focuses on evaluating and implementing state-of-the-art methods for long-tailed learning, as outlined in the survey *Deep Long-Tailed Learning: A Survey* by Zhang et al. The methodologies explored include advanced sampling strategies, re-weighted loss functions, and modifications to deep learning architectures tailored to imbalanced data.

A unique application of these methods is demonstrated on a custom dataset of moth images collected near the equator, where the goal is accurate species identification. Through a series of experiments, the thesis investigates how different approaches to long-tailed learning impact model performance across head, middle, and tail classes.

The findings contribute to understanding the efficacy of these methods and provide insights into best practices for handling real-world long-tailed datasets.

# Acknowledgements

I would like to thank my supervisor, Kim Bjerger, for their guidance, constructive feedback, and inspiring conversations throughout the development of this thesis. I am also deeply thankful to my family for providing me with room for my frustrations, and for cooking meals during the challenging times of this journey. I want to thank my friend, Line, for her valuable insights in writing a thesis. A special thanks to boyfriend for his outstanding patience and understanding.

# Contents

|  |            |
|--|------------|
| <b>Abstract</b>                                  | <b>iii</b> |
| <b>Acknowledgements</b>                          | <b>iv</b>  |
| <b>1 Introduction</b>                            | <b>1</b>   |
| 1.1 Problem Definition . . . . .                 | 1          |
| 1.1.1 Goals of this thesis . . . . .             | 1          |
| 1.1.2 Approach . . . . .                         | 1          |
| 1.1.3 Scope of this thesis . . . . .             | 2          |
| 1.2 Motivation . . . . .                         | 2          |
| 1.3 Reading Guide . . . . .                      | 2          |
| 1.4 Related Work . . . . .                       | 2          |
| <b>2 Background</b>                              | <b>3</b>   |
| 2.1 Long-Tailed Dataset . . . . .                | 3          |
| 2.2 Model Architectures . . . . .                | 3          |
| 2.2.1 Convolutional Neural Networks . . . . .    | 3          |
| 2.2.2 Visual Transformers . . . . .              | 3          |
| 2.3 Classic Long-Tailed Methods . . . . .        | 4          |
| 2.3.1 Class Re-balancing . . . . .               | 4          |
| 2.3.2 Information Augmentation . . . . .         | 6          |
| 2.3.3 Module Improvement . . . . .               | 6          |
| <b>3 Methodology</b>                             | <b>7</b>   |
| 3.1 Overview of Methodology Approach . . . . .   | 7          |
| 3.2 Algorithm Selection and Rationale . . . . .  | 7          |
| 3.3 Long-tailed Learning Techniques . . . . .    | 7          |
| 3.4 Loss Functions . . . . .                     | 7          |
| 3.5 Data Imbalance Handling Strategies . . . . . | 8          |
| 3.6 Evaluation Strategies . . . . .              | 8          |
| 3.7 Implementation Details . . . . .             | 8          |
| <b>4 Experimental Setup</b>                      | <b>9</b>   |
| 4.1 Dataset Specifications . . . . .             | 9          |
| 4.2 Data Preprocessing . . . . .                 | 9          |
| 4.3 Model Architecture Settings . . . . .        | 9          |

|          |  |           |
|----------|--|-----------|
| 4.4      | Training Configurations . . . . .                  | 9         |
| 4.5      | Evaluation Metrics . . . . .                       | 9         |
| 4.6      | Hardware and Software Configurations . . . . .     | 10        |
| 4.7      | Reproducibility Considerations . . . . .           | 10        |
| <b>5</b> | <b>Results and Analysis</b>                        | <b>11</b> |
| 5.1      | Overall Results . . . . .                          | 11        |
| 5.2      | Head, Middle, and Tail Class Performance . . . . . | 11        |
| 5.2.1    | MobileNetV2 . . . . .                              | 11        |
| 5.2.2    | ResNet50V2 . . . . .                               | 12        |
| 5.2.3    | ViT-B/16 . . . . .                                 | 13        |
| 5.2.4    | ConvNeXt Base . . . . .                            | 13        |
| 5.3      | Comparison of Loss Functions . . . . .             | 14        |
| 5.4      | Qualitative Results . . . . .                      | 14        |
| 5.5      | Summary and Discussion . . . . .                   | 15        |
| <b>6</b> | <b>Conclusion and Future Work</b>                  | <b>16</b> |
| 6.1      | Revisiting the Goals of the Thesis . . . . .       | 16        |
| 6.2      | Future Work . . . . .                              | 16        |
|          | <b>Bibliography</b>                                | <b>17</b> |
| <b>A</b> | <b>Results</b>                                     | <b>18</b> |
| A.1      | MobileNetV2 . . . . .                              | 18        |
| A.2      | ResNet50V2 . . . . .                               | 19        |
| A.3      | ViT-B/16 . . . . .                                 | 21        |
| A.4      | ConvNeXt Base . . . . .                            | 22        |

# Chapter 1

## Introduction

This thesis focuses on the problem with long-tailed datasets. The problem with training a deep learning model on long-tailed datasets is that the model will effectively the data from the classes with most samples, and not the classes with few samples. The finished model will then not recognize an input from the tail classes. Most real-world datasets follows a long-tailed structure, hence the need for a reliable method to detect examples of tail-class data. The aim of this thesis is to try out some of the methods tackling the long-tailed problem for deep learning described in the paper *Deep Long-Tailed Learning: A Survey* by Zhang et al.[1] to find a method for long-tailed learning that works on a specific long-tailed dataset of images of moths taken around equator. The goal of the moth dataset is to identify species.

### 1.1 Problem Definition

Define the problem formally, including key terms like "head classes" and "tail classes." Provide an example or visualization of a long-tailed dataset. Connect the problem definition to the moth dataset.

#### 1.1.1 Goals of this thesis

Outline the goals of the thesis, emphasizing optimizing performance on tail classes. Mention how these goals contribute to the field of long-tailed learning.

#### Hypothesis

#### 1.1.2 Approach

Summarize the approach to achieve goals, such as implementing and comparing methods from Zhang et al.'s survey.

### **1.1.3 Scope of this thesis**

Specify the scope: Focus is on image classification. Methods are tested on a specific dataset. The evaluation is limited to certain metrics.

## **1.2 Motivation**

Discuss why the problem is significant, including real-world implications. Mention the importance of biodiversity studies or the challenges of species identification with limited samples. Discuss broader impacts, such as how solving long-tailed learning problems can benefit other fields.

## **1.3 Reading Guide**

Mention what each chapter will cover and how they relate to each other.

## **1.4 Related Work**

A section that describes the work related to this thesis.



# Chapter 2

## Background

This chapter presents the different background topics of the thesis work, which are the long-tailed datasets, model architectures *Convolutional Neural Networks (CNN)* and *Visual Transformers (VT)*, the deep long-tailed learning methods *Class Re-balancing (CR)*, *Information Augmentation (IA)*, and *Module Improvement (MI)*. These topics will be explained for the reader.

Mention image classification, as it is the primary goal of this thesis.

### 2.1 Long-Tailed Dataset

A short introductory paragraph explaining why these topics are relevant and how they tie into the thesis. Contextualize the sections—e.g., "Long-tailed datasets are central to this thesis, as they represent the primary challenge. Model architectures and classic long-tailed methods form the foundation of the approaches explored in this work."

Mention the difference between class-imbalanced learning and long-tailed learning.

### 2.2 Model Architectures

Describe the role of deep learning models in handling long-tailed datasets.

#### 2.2.1 Convolutional Neural Networks

Add historical context. Mention specific CNNs used in this thesis (e.g., ResNet, MobileNet).

#### 2.2.2 Visual Transformers

Explain their advantages over CNNs for certain tasks. Mention why they are relevant for handling long-tailed datasets.

## 2.3 Classic Long-Tailed Methods

Introduce the three methods (CR, IA, MI) with a brief explanation of their purpose.

Following the paper *Deep Long-Tailed Learning: A Survey* [1], the existing deep long-tailed learning methods are grouped into three main categories based on their technical approach: class re-balancing, information augmentation, and module improvement. These categories are further divided onto sub-categories: re-sampling, class-sensitive learning, logit adjustment, transfer learning, data augmentation, representation learning, classifier desing, decoupled training, and ensemble learning (TODO: create a figure like figure 2 in the paper). This thesis does not aim to examine all the beforementioned method, but aims to find a deep learning approach to a specific problem. The backgrounds of the methods used in this thesis are described in this section.

### 2.3.1 Class Re-balancing

The class re-balancing method aims to re-balance the effect of the imbalanced training dataset, and has three main sub-categories: re-sampling, class-sensitive learning, and logit adjustment [1].

#### Re-sampling

The traditional way to sample when training deep networks is bases on mini-batch gradient descent with random sampling. This means that each sample has an equal probability of being sampled. When sampling from an imbalanced dataset, samples from head classes naturally occur more often, and thus have higher chance of being sampled than samples from tail classes, making the resulting deep models biased towards head classes. Re-sampling is a method that adresses this problem by adjusting the number of samples per class in each sample batch for model training.

#### Class-sensitive Learning

Class-sensitive learning incorporates strategies to adjust the loss function, making it more sensitive to the imbalanced nature of the dataset. This approach directly modifies the optimization process to prioritize learning from under-represented tail classes.

TODO: Mention re-weighting and re-margining. Table/overview of loss functions as in the paper.

#### Loss Functions for Class-Sensitive Learning

The loss function serves as a measure of the model's fitness to the data, quantifying the distance between the actual and predicted values of the target. Typically, the

loss is represented as a nonnegative value, where smaller values indicate a better fit, and a perfect fit corresponds to a loss of zero [2].

Conventional training of deep networks using the softmax cross-entropy loss often overlooks class imbalance. This results in uneven gradients for different classes, leading to suboptimal performance on underrepresented classes. To mitigate this issue, modifications to the loss function are introduced to ensure a more balanced contribution from each class during training. One such technique is re-weighting which adjusts the training loss for different classes by assigning a specific weight to each class [1]. The softmax cross-entropy loss is used as a baseline, and is described below along with the loss functions for re-weighting.

**Softmax Cross-Entropy Loss** The *Softmax-Cross-Entropy loss*, often referred to as *softmax loss*, is a widely used combination for training deep neural networks in classification tasks, including image classification. It is particularly effective for multi-class problems, where the goal is to assign an input image to one of several predefined categories [3] [4].

The *Softmax* function transforms the raw output scores (logits) of the final layer of a neural network into a probability distribution over  $K$  classes. For an input  $\mathbf{z} = [z_1, z_2, \dots, z_K]$ , the Softmax function for class  $i$  is defined as:

$$P(y = i \mid \mathbf{z}) = \frac{\exp(z_i)}{\sum_{j=1}^K \exp(z_j)} \quad (2.1)$$

Here,  $\exp(z_i)$  ensures that all values are positive, and dividing by the sum normalizes the probabilities so that they sum to 1. This normalization is crucial for classification, as it allows the network's outputs to represent the likelihood of each class.

The *Cross-Entropy loss* measures the difference between the predicted probability distribution  $\mathbf{P}$  (produced by Softmax) and the true distribution  $\mathbf{y}$  (the one-hot encoded ground truth). It is defined as:

$$\mathcal{L}_{\text{CE}} = - \sum_{i=1}^K y_i \log(P(y = i \mid \mathbf{z})) \quad (2.2)$$

For a single example where the true class is  $c$ , this simplifies to:

$$\mathcal{L}_{\text{CE}} = - \log(P(y = c \mid \mathbf{z})) \quad (2.3)$$

This formulation penalizes incorrect predictions by heavily weighting the log of the predicted probability for the true class. The loss is minimized when the predicted probability  $P(y = c \mid \mathbf{z})$  approaches 1, indicating high confidence in the correct class.

This combination has become the de facto standard for image classification tasks, providing a robust and mathematically sound framework for training deep neural networks.

**Weighted Softmax Cross-Entropy Loss** The *Weighted Softmax Cross-Entropy loss*, often referred to as *weighted softmax loss*, is a variant of the standard

softmax cross-entropy loss, designed to address imbalanced datasets [4] [5]. By assigning different weights to each class, this method ensures that underrepresented classes contribute more to the overall loss, improving the model’s performance on minority classes. The weighted cross-entropy loss applies class-specific weights to the standard cross-entropy formulation. It is defined as:

$$\mathcal{L}_{\text{WCE}} = - \sum_{i=1}^K w_i y_i \log(P(y = i | \mathbf{z})) \quad (2.4)$$

Where  $w_i$  is the weight for class  $i$ , reflecting its relative importance,  $y_i$  is the one-hot encoded true label for class  $i$ , and  $P(y = i | \mathbf{z})$  is the predicted probability for class  $i$ .

For a single example where the true class is  $c$ , the loss simplifies to:

$$\mathcal{L}_{\text{WCE}} = -w_c \log(P(y = c | \mathbf{z})) \quad (2.5)$$

This weighted formulation ensures that minority classes contribute more to the overall loss, addressing the imbalance during training and improving the model’s performance on underrepresented classes.

**Focal Loss** Focal Loss, introduced by Lin et al. (2017) [5], addresses the challenges of extreme class imbalance in classification tasks by dynamically scaling the standard cross-entropy loss. Focal Loss mitigates the issue of imbalanced datasets by down-weighting the loss contributions from well-classified examples and focusing on misclassified examples during training.

**Class-Balanced Loss** *Class-balanced loss*, introduced by Cui et al. (2019) [6], ...

**Balanced Softmax Loss** *Balanced Softmax loss*, introduced by Ren et al. (2020) [7], ...

**LDAM Loss** *LDAM loss*, introduced by Cao et al. (2019) [8], ...

**Equalization Loss** *Equalization loss*, introduced by Tan et al. (2020) [9], ...

## 2.3.2 Information Augmentation

Data augmentation techniques tailored for long-tailed datasets.

### Transfer Learning

### Data Augmentation

## 2.3.3 Module Improvement

Architectural changes to improve tail-class representation.

# Chapter 3

## Methodology

This chapter describes the methods and approaches used in the experiments. This includes dataset preparation, models, loss functions, etc.

### 3.1 Overview of Methodology Approach

An overall description of the approach to tackling the long-tailed dataset problem, including an explanation of the strategy, such as balancing techniques and model selection.

### 3.2 Algorithm Selection and Rationale

Description of the model architectures chosen, and why they are appropriate for long-tailed learning. Discussion of the strengths and limitations of these models in addressing the challenges posed by imbalanced data.

### 3.3 Long-tailed Learning Techniques

Description of the specific methods used to address class imbalance, such as data sampling, class re-weighting, etc. Justification for selecting these techniques, potentially referencing prior research (from Deep Long-Tailed Learning: A Survey by Zhang et al.).

### 3.4 Loss Functions

Explanation of the different loss functions explored, such as cross-entropy, focal loss, LDAM loss, etc., and their relevance for long-tailed learning. Rationale for each loss function's inclusion, focusing on its expected benefits for imbalanced classes and how it addresses the bias toward majority classes.

### **3.5 Data Imbalance Handling Strategies**

Detailed explanation of the techniques for creating and handling an imbalanced dataset, such as generating imbalanced training and test sets.

### **3.6 Evaluation Strategies**

Justification for the metrics and evaluation approach, such as using weighted or macro F1 scores. Explanation of how the performance is assessed across different class groups to capture the model's performance on minority classes.

### **3.7 Implementation Details**

Technical explanations of any unique or customized methods implemented in code, for example the custom dataset.

# Chapter 4

## Experimental Setup

This chapter focuses on the on the implementation details of the experiments conducted in this thesis. Here, the specifics of the training configurations are described.

### 4.1 Dataset Specifications

Details about the dataset(s) used, including size, source, and preprocessing steps. Description of class imbalance characteristics and the train/validation/test splits.

### 4.2 Data Preprocessing

Any transformations, augmentations, or normalization applied to the dataset before feeding it to the model. Information on how the class imbalance is handled (re-sampling or synthetic data generation).

### 4.3 Model Architecture Settings

Description of the models used, including any specific architecture choices, hyperparameters, or modifications. Brief details on why these models were chosen.

### 4.4 Training Configurations

Hyperparameters, such as batch size, learning rate, optimizer type, and regularization techniques (dropout, weight decay). Any specific settings for handling long-tailed data, such as DRW.

### 4.5 Evaluation Metrics

Explanation of the metrics used to assess model performance. Justification for choosing each metric.

## 4.6 Hardware and Software Configurations

Hardware details. Software environment, including the versions of libraries and frameworks.

## 4.7 Reproducibility Considerations

Steps taken to ensure that results can be reproduced, such as random seed initialization and details on dataset versions. Scripts, configurations, or instructions for reproducing experiments.



# Chapter 5

## Results and Analysis

Presentation of the results, with tables, charts, and explanations for each tested method's performance.

Brief overview of the chapter's purpose. Recap the evaluation goals (model performance across head, middle, and tail classes, and comparing methods).

### 5.1 Overall Results

Present the performance of all tested models and methods. Use tables or charts to summarize key results. Highlight trends or notable observations across the methods.

### 5.2 Head, Middle, and Tail Class Performance

Break down the performance into head, middle, and tail class groups. Include visualizations. Discuss how well the methods balance performance across these groups, particularly focusing on tail classes.

#### 5.2.1 MobileNetV2

Table 5.1 show the top 1 accuracies for MobileNetV2 on various loss functions.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax               | 0.7978   | 0.8059      | 0.8069 | 0.7870 | 0.8684 |
| Focal loss            | 0.8014   | 0.8011      | 0.7998 | 0.7870 | 0.8947 |
| Weighted Softmax loss | 0.7978   | 0.8059      | 0.8069 | 0.7870 | 0.8684 |
| Class-balanced loss   | 0.7978   | 0.8059      | 0.8069 | 0.7870 | 0.8684 |
| Balanced Softmax loss | 0.8034   | 0.8030      | 0.8069 | 0.7574 | 0.9211 |
| Equalization loss     | 0.7994   | 0.8040      | 0.8057 | 0.7692 | 0.9211 |
| LDAM loss             | 0.7828   | 0.7821      | 0.7808 | 0.7574 | 0.9211 |

Table 5.1: Evaluation results for MobileNetV2 trained on the custom balanced dataset, showing Acc1.

Table 5.2 show the top 1 accuracies for MobileNetV2 on various loss functions.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax               | 0.5282   | 0.7735      | 0.8341 | 0.5917 | 0.2368 |
| Focal loss            | 0.5200   | 0.7745      | 0.8389 | 0.5917 | 0.1579 |
| Weighted Softmax loss | 0.5016   | 0.7231      | 0.7808 | 0.5503 | 0.2105 |
| Class-balanced loss   | 0.1936   | 0.0913      | 0.0521 | 0.2485 | 0.2632 |
| Balanced Softmax loss | 0.5796   | 0.7650      | 0.8069 | 0.6331 | 0.4211 |
| Equalization loss     | 0.5310   | 0.7650      | 0.8235 | 0.5917 | 0.2368 |
| LDAM loss             | 0.4264   | 0.5899      | 0.6137 | 0.5444 | 0.2632 |

Table 5.2: Evaluation results for MobileNetV2 trained on the long-tailed dataset showing Acc1.

## 5.2.2 ResNet50V2

Table 5.3 show the top 1 accuracies for ResNet50V2 on various loss functions.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax loss          | 0.8324   | 0.8421      | 0.8448 | 0.8047 | 0.9474 |
| Focal loss            | 0.8310   | 0.8344      | 0.8341 | 0.8166 | 0.9211 |
| Weighted Softmax loss | 0.8324   | 0.8421      | 0.8448 | 0.8047 | 0.9474 |
| Class-balanced loss   | 0.8324   | 0.8421      | 0.8448 | 0.8047 | 0.9474 |
| Balanced Softmax loss | 0.8310   | 0.8430      | 0.8460 | 0.8107 | 0.9211 |
| Equalization loss     | 0.8292   | 0.8373      | 0.8412 | 0.7929 | 0.9474 |
| LDAM loss             | 0.7990   | 0.7983      | 0.8069 | 0.7337 | 0.8947 |

Table 5.3: Evaluation results for ResNet50V2 trained on the custom balanced dataset, showing Acc1.

Table 5.4 show the top 1 accuracies for ResNet50V2 on various loss functions.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax loss          | 0.5522   | 0.7954      | 0.8531 | 0.6391 | 0.2105 |
| Focal loss            | 0.5456   | 0.7935      | 0.8483 | 0.6272 | 0.3158 |
| Weighted Softmax loss | 0.4976   | 0.7336      | 0.7915 | 0.5562 | 0.2368 |
| Class-balanced loss   | 0.2052   | 0.1836      | 0.1445 | 0.3787 | 0.1842 |
| Balanced Softmax loss | 0.5908   | 0.7916      | 0.8270 | 0.6568 | 0.6053 |
| Equalization loss     | 0.5452   | 0.7897      | 0.8389 | 0.6450 | 0.3421 |
| LDAM loss             | 0.3742   | 0.5937      | 0.6469 | 0.4438 | 0.0789 |

Table 5.4: Evaluation results for ResNet50V2 trained on the long-tailed dataset, showing Acc1.

### 5.2.3 ViT-B/16

Table 5.5 show the top 1 accuracies for ViT-B/16 on various loss functions.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax loss          | 0.5620   | 0.5671      | 0.5521 | 0.6036 | 0.7368 |
| Focal loss            | 0.5516   | 0.5538      | 0.5438 | 0.5680 | 0.7105 |
| Weighted Softmax loss | 0.5620   | 0.5671      | 0.5521 | 0.6036 | 0.7368 |
| Class-balanced loss   | 0.5620   | 0.5671      | 0.5521 | 0.6036 | 0.7368 |
| Balanced Softmax loss | 0.5628   | 0.5642      | 0.5640 | 0.5325 | 0.7105 |
| Equalization loss     | 0.5634   | 0.5519      | 0.5462 | 0.5503 | 0.6842 |
| LDAM loss             | 0.5906   | 0.6013      | 0.5924 | 0.6095 | 0.7632 |

Table 5.5: Evaluation results for ViT-B/16 trained on the custom balanced dataset, showing Acc1.

Table 5.6 show the top 1 accuracies for ViT-B/16 on various loss functions.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax loss          | 0.2254   | 0.4367      | 0.5071 | 0.1775 | 0.0263 |
| Focal loss            | 0.2210   | 0.4206      | 0.4834 | 0.1953 | 0.0263 |
| Weighted Softmax loss | 0.1284   | 0.1760      | 0.1919 | 0.1302 | 0.0263 |
| Class-balanced loss   | 0.0558   | 0.0076      | 0.0000 | 0.0237 | 0.1053 |
| Balanced Softmax loss | 0.2460   | 0.4244      | 0.4822 | 0.2130 | 0.0789 |
| Equalization loss     | 0.2168   | 0.4215      | 0.4893 | 0.1716 | 0.0263 |
| LDAM loss             | 0.5906   | 0.6013      | 0.5924 | 0.6095 | 0.7632 |

Table 5.6: Evaluation results for ViT-B/16 trained on the long-tailed dataset, showing Acc1.

### 5.2.4 ConvNeXt Base

Table 5.7 show the top 1 accuracies for ConvNeXt Base on various loss functions.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax loss          | 0.8332   | 0.8535      | 0.8566 | 0.8166 | 0.9474 |
| Focal loss            | 0.8314   | 0.8487      | 0.8507 | 0.8284 | 0.8947 |
| Weighted Softmax loss | 0.8332   | 0.8535      | 0.8566 | 0.8166 | 0.9474 |
| Class-balanced loss   | 0.8332   | 0.8535      | 0.8566 | 0.8166 | 0.9474 |
| Balanced Softmax loss | 0.8364   | 0.8344      | 0.8365 | 0.7988 | 0.9474 |
| Equalization loss     | 0.8318   | 0.8468      | 0.8448 | 0.8343 | 0.9474 |
| LDAM loss             | data     | data        | data   | data   | data   |

Table 5.7: Evaluation results for ConvNeXt Base trained on the custom balanced dataset, showing Acc1.

Table 5.8 show the top 1 accuracies for ConvNeXt Base on various loss functions.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax loss          | 0.5972   | 0.8316      | 0.8898 | 0.6568 | 0.3158 |
| Focal loss            | 0.5938   | 0.8145      | 0.8685 | 0.6568 | 0.3158 |
| Weighted Softmax loss | 0.4090   | 0.6356      | 0.6848 | 0.4911 | 0.1842 |
| Class-balanced loss   | 0.0142   | 0.0019      | 0.0000 | 0.0000 | 0.0526 |
| Balanced Softmax loss | 0.6460   | 0.8230      | 0.8685 | 0.6509 | 0.5789 |
| Equalization loss     | 0.5956   | 0.8278      | 0.8768 | 0.6923 | 0.3421 |
| LDAM loss             | data     | data        | data   | data   | data   |

Table 5.8: Evaluation results for ConvNeXt Base trained on the long-tailed dataset, showing Acc1.

### 5.3 Comparison of Loss Functions

Analyze how different loss functions impact performance. Use visualizations to compare results ( per-class performance or confusion matrices). Discuss strengths and weaknesses of each loss function.

### 5.4 Qualitative Results

Optional.

Provide examples of correctly and incorrectly classified samples, especially for tail classes. Include visualizations or images of difficult cases to highlight challenges in tail-class prediction.

## 5.5 Summary and Discussion

Recap the key findings, such as which methods or loss functions performed best and why. Connect these findings to the thesis objectives and broader implications for long-tailed learning.

# Chapter 6

## Conclusion and Future Work

Summary of the work, contributions, and suggestions for future improvements or research directions.

### 6.1 Revisiting the Goals of the Thesis

### 6.2 Future Work

# Bibliography

- [1] Yifan Zhang et al. “Deep long-tailed learning: A survey”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* (2023).
- [2] Aston Zhang et al. *Dive into Deep Learning*. <https://D2L.ai>. Cambridge University Press, 2023.
- [3] X. L. Chaitanya Asawa. *CS231n: Convolutional Neural Networks for Visual Recognition*. <http://cs231n.github.io/>. Stanford, [Online]. 2024.
- [4] *torch.nn.CrossEntropyLoss* — *PyTorch documentation*. <https://pytorch.org/docs/stable/generated/torch.nn.CrossEntropyLoss.html>. Accessed: 2024-11-20.
- [5] Tsung-Yi Lin et al. *Focal Loss for Dense Object Detection*. 2018. arXiv: 1708.02002 [cs.CV]. URL: <https://arxiv.org/abs/1708.02002>.
- [6] Yin Cui et al. *Class-Balanced Loss Based on Effective Number of Samples*. 2019. arXiv: 1901.05555 [cs.CV]. URL: <https://arxiv.org/abs/1901.05555>.
- [7] Jiawei Ren et al. *Balanced Meta-Softmax for Long-Tailed Visual Recognition*. 2020. arXiv: 2007.10740 [cs.LG]. URL: <https://arxiv.org/abs/2007.10740>.
- [8] Kaidi Cao et al. *Learning Imbalanced Datasets with Label-Distribution-Aware Margin Loss*. 2019. arXiv: 1906.07413 [cs.LG]. URL: <https://arxiv.org/abs/1906.07413>.
- [9] Jingru Tan et al. *Equalization Loss for Long-Tailed Object Recognition*. 2020. arXiv: 2003.05176 [cs.CV]. URL: <https://arxiv.org/abs/2003.05176>.

# Appendix A

## Results

Tables of the results from training.

### A.1 MobileNetV2

*MobileNetV2 trained on custom balanced dataset and imbalanced dataset on different loss functions.*

Table A.1 show the top 1 accuracies for MobileNetV2 on various loss functions. Table A.2 show the loss, top 1 accuracy, and F1 score. Table A.3 show the top 1 accuracies for MobileNetV2 on various loss functions. Table A.4 show the loss, top 1 accuracy, and F1 score.

| Loss Function         | Balanced | Long-tailed | Head     | Middle | Tail   |
|-----------------------|----------|-------------|----------|--------|--------|
| Softmax               | 0.7978   | 0.8059      | 0.8069   | 0.7870 | 0.8684 |
| Focal loss            | 0.8014   | 0.8011      | 0.7998 4 | 0.7870 | 0.8947 |
| Weighted Softmax loss | 0.7978   | 0.8059      | 0.8069   | 0.7870 | 0.8684 |
| Class-balanced loss   | 0.7978   | 0.8059      | 0.8069   | 0.7870 | 0.8684 |
| Balanced Softmax loss | 0.8034   | 0.8030      | 0.8069   | 0.7574 | 0.9211 |
| Equalization loss     | 0.7994   | 0.8040      | 0.8057   | 0.7692 | 0.9211 |
| LDAM loss             | 0.7828   | 0.7821      | 0.7808   | 0.7574 | 0.9211 |

Table A.1: Evaluation results for MobileNetV2 trained on the custom balanced dataset, showing Acc1.



| Loss Function    | Balanced |        |        | Long-tailed |        |        | Head    |        |        | Middle  |        |        | Tail   |        |        |
|------------------|----------|--------|--------|-------------|--------|--------|---------|--------|--------|---------|--------|--------|--------|--------|--------|
|                  | Loss     | Acc1   | F1     | Loss        | Acc1   | F1     | Loss    | Acc1   | F1     | Loss    | Acc1   | F1     | Loss   | Acc1   | F1     |
| Softmax          | 1.1455   | 0.7978 | 0.7967 | 1.1415      | 0.8059 | 0.8208 | 1.1208  | 0.8069 | 0.8587 | 1.3060  | 0.7870 | 0.8368 | 0.8690 | 0.8684 | 0.8684 |
| Focal Loss       | 0.6765   | 0.8014 | 0.8001 | 0.7063      | 0.8011 | 0.8175 | 0.7005  | 0.7998 | 0.8531 | 0.8028  | 0.7870 | 0.8293 | 0.4055 | 0.8947 | 0.8860 |
| Weighted Softmax | 1.1455   | 0.7978 | 0.7967 | 1.1415      | 0.8059 | 0.8208 | 1.1208  | 0.8069 | 0.8587 | 1.3060  | 0.7870 | 0.8368 | 0.8690 | 0.8684 | 0.8684 |
| Class-balanced   | 1.1455   | 0.7978 | 0.7967 | 1.1415      | 0.8059 | 0.8208 | 1.1208  | 0.8069 | 0.8587 | 1.3060  | 0.7870 | 0.8368 | 0.8690 | 0.8684 | 0.8684 |
| Balanced Softmax | 1.1289   | 0.8034 | 0.8011 | 1.1848      | 0.8030 | 0.8145 | 1.1469  | 0.8069 | 0.8553 | 1.4407  | 0.7574 | 0.8123 | 0.8872 | 0.9211 | 0.9298 |
| Equalization     | 0.9992   | 0.7994 | 0.7983 | 1.0385      | 0.8040 | 0.8192 | 1.0118  | 0.8057 | 0.8564 | 1.2539  | 0.7692 | 0.8213 | 0.6035 | 0.9211 | 0.9211 |
| LDAM             | 13.8126  | 0.7828 | 0.7817 | 13.5566     | 0.7821 | 0.7955 | 13.6884 | 0.7808 | 0.8325 | 14.7496 | 0.7574 | 0.8016 | 5.3231 | 0.9211 | 0.9035 |

Table A.2: Evaluation results for MobileNetV2 trained on the custom balanced dataset, showing Loss, Acc1, and F1 scores for each dataset split.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax               | 0.5282   | 0.7735      | 0.8341 | 0.5917 | 0.2368 |
| Focal loss            | 0.5200   | 0.7745      | 0.8389 | 0.5917 | 0.1579 |
| Weighted Softmax loss | 0.5016   | 0.7231      | 0.7808 | 0.5503 | 0.2105 |
| Class-balanced loss   | 0.1936   | 0.0913      | 0.0521 | 0.2485 | 0.2632 |
| Balanced Softmax loss | 0.5796   | 0.7650      | 0.8069 | 0.6331 | 0.4211 |
| Equalization loss     | 0.5310   | 0.7650      | 0.8235 | 0.5917 | 0.2368 |
| LDAM loss             | 0.4264   | 0.5899      | 0.6137 | 0.5444 | 0.2632 |

Table A.3: Evaluation results for MobileNetV2 trained on the long-tailed dataset showing Acc1.

| Loss Function    | Balanced |        |        | Long-tailed |        |        | Head   |        |        | Middle  |        |        | Tail    |        |        |
|------------------|----------|--------|--------|-------------|--------|--------|--------|--------|--------|---------|--------|--------|---------|--------|--------|
|                  | Loss     | Acc1   | F1     | Loss        | Acc1   | F1     | Loss   | Acc1   | F1     | Loss    | Acc1   | F1     | Loss    | Acc1   | F1     |
| Softmax          | 3.2503   | 0.5282 | 0.4884 | 1.2212      | 0.7735 | 0.7578 | 0.8136 | 0.8341 | 0.8492 | 2.4604  | 0.5917 | 0.6444 | 4.7629  | 0.2368 | 0.2544 |
| Focal Loss       | 2.3526   | 0.5200 | 0.4818 | 0.8022      | 0.7745 | 0.7602 | 0.5177 | 0.8389 | 0.8528 | 1.5864  | 0.5917 | 0.6625 | 3.6343  | 0.1579 | 0.1667 |
| Weighted Softmax | 3.1412   | 0.5016 | 0.4690 | 1.2817      | 0.7231 | 0.7104 | 0.8786 | 0.7808 | 0.8015 | 2.3365  | 0.5503 | 0.6213 | 5.1836  | 0.2105 | 0.1912 |
| Class-balanced   | 4.3308   | 0.1936 | 0.1751 | 4.2181      | 0.0913 | 0.0854 | 4.4197 | 0.0521 | 0.0795 | 2.8788  | 0.2485 | 0.2767 | 6.3575  | 0.2632 | 0.2368 |
| Balanced Softmax | 3.1185   | 0.5796 | 0.5572 | 1.1630      | 0.7650 | 0.7685 | 0.7989 | 0.8069 | 0.8422 | 2.1612  | 0.6331 | 0.6872 | 4.8108  | 0.4211 | 0.4123 |
| Equalization     | 3.0593   | 0.5310 | 0.4911 | 1.1563      | 0.7650 | 0.7499 | 0.7241 | 0.8235 | 0.8398 | 2.4487  | 0.5917 | 0.6524 | 4.9284  | 0.2368 | 0.2544 |
| LDAM             | 21.4896  | 0.4264 | 0.3980 | 7.9893      | 0.5899 | 0.5909 | 5.6756 | 0.6137 | 0.6581 | 10.3379 | 0.5444 | 0.6121 | 49.1197 | 0.2632 | 0.2895 |

Table A.4: Evaluation results for MobileNetV2 trained on the long-tailed dataset, showing Loss, Acc1, and F1 scores for each dataset split.

## A.2 ResNet50V2

*ResNet50V2 trained on custom balanced dataset and imbalanced dataset on different loss functions.*

Table A.5 show the top 1 accuracies for ResNet50V2 on various loss functions. Table A.6 show the loss, top 1 accuracy, and F1 score.

Table A.7 show the top 1 accuracies for ResNet50V2 on various loss functions. Table A.8 show the loss, top 1 accuracy, and F1 score.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax loss          | 0.8324   | 0.8421      | 0.8448 | 0.8047 | 0.9474 |
| Focal loss            | 0.8310   | 0.8344      | 0.8341 | 0.8166 | 0.9211 |
| Weighted Softmax loss | 0.8324   | 0.8421      | 0.8448 | 0.8047 | 0.9474 |
| Class-balanced loss   | 0.8324   | 0.8421      | 0.8448 | 0.8047 | 0.9474 |
| Balanced Softmax loss | 0.8310   | 0.8430      | 0.8460 | 0.8107 | 0.9211 |
| Equalization loss     | 0.8292   | 0.8373      | 0.8412 | 0.7929 | 0.9474 |
| LDAM loss             | 0.7990   | 0.7983      | 0.8069 | 0.7337 | 0.8947 |

Table A.5: Evaluation results for ResNet50V2 trained on the custom balanced dataset, showing Acc1.

| Loss Function    | Balanced |        |        | Long-tailed |        |        | Head   |        |        | Middle  |        |        | Tail   |        |        |
|------------------|----------|--------|--------|-------------|--------|--------|--------|--------|--------|---------|--------|--------|--------|--------|--------|
|                  | Loss     | Acc1   | F1     | Loss        | Acc1   | F1     | Loss   | Acc1   | F1     | Loss    | Acc1   | F1     | Loss   | Acc1   | F1     |
| Softmax          | 0.9823   | 0.8324 | 0.8310 | 0.9874      | 0.8421 | 0.8520 | 0.9917 | 0.8448 | 0.8860 | 1.0934  | 0.8047 | 0.8467 | 0.4205 | 0.9474 | 0.9386 |
| Focal Loss       | 0.5627   | 0.8310 | 0.8300 | 0.5578      | 0.8344 | 0.8474 | 0.5555 | 0.8341 | 0.8788 | 0.6294  | 0.8166 | 0.8607 | 0.2920 | 0.9211 | 0.9123 |
| Weighted Softmax | 0.9823   | 0.8324 | 0.8310 | 0.9874      | 0.8421 | 0.8520 | 0.9917 | 0.8448 | 0.8860 | 1.0934  | 0.8047 | 0.8467 | 0.4205 | 0.9474 | 0.9386 |
| Class-balanced   | 0.9823   | 0.8324 | 0.8310 | 0.9874      | 0.8421 | 0.8520 | 0.9917 | 0.8448 | 0.8860 | 1.0934  | 0.8047 | 0.8467 | 0.4205 | 0.9474 | 0.9386 |
| Balanced Softmax | 1.0198   | 0.8310 | 0.8301 | 0.9689      | 0.8430 | 0.8549 | 0.9601 | 0.8460 | 0.8893 | 1.1309  | 0.8107 | 0.8539 | 0.4440 | 0.9211 | 0.9123 |
| Equalization     | 0.8795   | 0.8292 | 0.8279 | 0.9079      | 0.8373 | 0.8495 | 0.8888 | 0.8412 | 0.8877 | 1.1374  | 0.7929 | 0.8453 | 0.2495 | 0.9474 | 0.9386 |
| LDAM             | 9.8339   | 0.7990 | 0.7979 | 10.1092     | 0.7983 | 0.8119 | 9.8723 | 0.8069 | 0.8596 | 12.5229 | 0.7337 | 0.7823 | 4.6362 | 0.8947 | 0.8772 |

Table A.6: Evaluation results for ResNet50V2 trained on the custom balanced dataset, showing Loss, Acc1, and F1 scores for each dataset split.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax loss          | 0.5522   | 0.7954      | 0.8531 | 0.6391 | 0.2105 |
| Focal loss            | 0.5456   | 0.7935      | 0.8483 | 0.6272 | 0.3158 |
| Weighted Softmax loss | 0.4976   | 0.7336      | 0.7915 | 0.5562 | 0.2368 |
| Class-balanced loss   | 0.2052   | 0.1836      | 0.1445 | 0.3787 | 0.1842 |
| Balanced Softmax loss | 0.5908   | 0.7916      | 0.8270 | 0.6568 | 0.6053 |
| Equalization loss     | 0.5452   | 0.7897      | 0.8389 | 0.6450 | 0.3421 |
| LDAM loss             | 0.3742   | 0.5937      | 0.6469 | 0.4438 | 0.0789 |

Table A.7: Evaluation results for ResNet50V2 trained on the long-tailed dataset, showing Acc1.

| Loss Function    | Balanced |        |        | Long-tailed |        |        | Head   |        |        | Middle  |        |        | Tail    |        |        |
|------------------|----------|--------|--------|-------------|--------|--------|--------|--------|--------|---------|--------|--------|---------|--------|--------|
|                  | Loss     | Acc1   | F1     | Loss        | Acc1   | F1     | Loss   | Acc1   | F1     | Loss    | Acc1   | F1     | Loss    | Acc1   | F1     |
| Softmax          | 3.0907   | 0.5522 | 0.5138 | 1.0524      | 0.7954 | 0.7798 | 0.6888 | 0.8531 | 0.8654 | 2.0330  | 0.6391 | 0.6996 | 4.7658  | 0.2105 | 0.2018 |
| Focal Loss       | 2.0718   | 0.5456 | 0.5089 | 0.6284      | 0.7935 | 0.7789 | 0.3983 | 0.8483 | 0.8583 | 1.3258  | 0.6272 | 0.7054 | 2.6364  | 0.3158 | 0.3158 |
| Weighted Softmax | 3.7904   | 0.4976 | 0.4591 | 1.3481      | 0.7336 | 0.7198 | 0.8630 | 0.7915 | 0.8098 | 2.2625  | 0.5562 | 0.6209 | 6.9808  | 0.2368 | 0.2456 |
| Class-balanced   | 4.5887   | 0.2052 | 0.1928 | 3.7422      | 0.1836 | 0.1932 | 3.7880 | 0.1445 | 0.2045 | 2.4884  | 0.3787 | 0.4138 | 8.3052  | 0.1842 | 0.1737 |
| Balanced Softmax | 3.1081   | 0.5908 | 0.5654 | 1.0452      | 0.7916 | 0.7895 | 0.6873 | 0.8270 | 0.8602 | 2.3422  | 0.6568 | 0.7135 | 3.2275  | 0.6053 | 0.5965 |
| Equalization     | 3.0166   | 0.5452 | 0.5071 | 1.0315      | 0.7897 | 0.7756 | 0.7418 | 0.8389 | 0.8511 | 1.8754  | 0.6450 | 0.7061 | 3.6342  | 0.3421 | 0.3509 |
| LDAM             | 22.7933  | 0.3742 | 0.3337 | 8.2056      | 0.5937 | 0.5784 | 5.3320 | 0.6469 | 0.6680 | 12.3074 | 0.4438 | 0.5450 | 53.4080 | 0.0789 | 0.0789 |

Table A.8: Evaluation results for ResNet50V2 trained on the long-tailed dataset, showing Loss, Acc1, and F1 scores for each dataset split.

### A.3 ViT-B/16

*ViT-B/16 trained on custom balanced dataset and imbalanced dataset on different loss functions.*

Table A.9 show the top 1 accuracies for ViT-B/16 on various loss functions. Table A.10 show the loss, top 1 accuracy, and F1 score.

Table A.11 show the top 1 accuracies for ViT-B/16 on various loss functions. Table A.12 show the loss, top 1 accuracy, and F1 score.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax loss          | 0.5620   | 0.5671      | 0.5521 | 0.6036 | 0.7368 |
| Focal loss            | 0.5516   | 0.5538      | 0.5438 | 0.5680 | 0.7105 |
| Weighted Softmax loss | 0.5620   | 0.5671      | 0.5521 | 0.6036 | 0.7368 |
| Class-balanced loss   | 0.5620   | 0.5671      | 0.5521 | 0.6036 | 0.7368 |
| Balanced Softmax loss | 0.5628   | 0.5642      | 0.5640 | 0.5325 | 0.7105 |
| Equalization loss     | 0.5634   | 0.5519      | 0.5462 | 0.5503 | 0.6842 |
| LDAM loss             | 0.5906   | 0.6013      | 0.5924 | 0.6095 | 0.7632 |

Table A.9: Evaluation results for ViT-B/16 trained on the custom balanced dataset, showing Acc1.

| Loss Function    | Balanced |        |        | Long-tailed |        |        | Head    |        |        | Middle  |        |        | Tail    |        |        |
|------------------|----------|--------|--------|-------------|--------|--------|---------|--------|--------|---------|--------|--------|---------|--------|--------|
|                  | Loss     | Acc1   | F1     | Loss        | Acc1   | F1     | Loss    | Acc1   | F1     | Loss    | Acc1   | F1     | Loss    | Acc1   | F1     |
| Softmax          | 4.6431   | 0.5620 | 0.5593 | 4.6089      | 0.5671 | 0.5951 | 4.6420  | 0.5521 | 0.6367 | 4.7263  | 0.6036 | 0.6648 | 3.3521  | 0.7368 | 0.7281 |
| Focal Loss       | 2.3473   | 0.5516 | 0.5488 | 2.3562      | 0.5538 | 0.5869 | 2.4288  | 0.5438 | 0.6324 | 2.2330  | 0.5680 | 0.6355 | 1.2929  | 0.7105 | 0.6930 |
| Weighted Softmax | 4.6431   | 0.5620 | 0.5593 | 4.6089      | 0.5671 | 0.5951 | 4.6420  | 0.5521 | 0.6367 | 4.7263  | 0.6036 | 0.6648 | 3.3521  | 0.7368 | 0.7281 |
| Class-balanced   | 4.6431   | 0.5620 | 0.5593 | 4.6089      | 0.5671 | 0.5951 | 4.6420  | 0.5521 | 0.6367 | 4.7263  | 0.6036 | 0.6648 | 3.3521  | 0.7368 | 0.7281 |
| Balanced Softmax | 4.7131   | 0.5628 | 0.5592 | 4.6809      | 0.5642 | 0.5929 | 4.7739  | 0.5640 | 0.6471 | 4.8161  | 0.5325 | 0.5998 | 2.0138  | 0.7105 | 0.7105 |
| Equalization     | 4.2603   | 0.5634 | 0.5614 | 4.4906      | 0.5519 | 0.5884 | 4.6109  | 0.5462 | 0.6410 | 4.3952  | 0.5503 | 0.6014 | 2.0079  | 0.6842 | 0.6754 |
| LDAM             | 48.2745  | 0.5906 | 0.5926 | 47.4149     | 0.6013 | 0.6348 | 49.6692 | 0.5924 | 0.6790 | 42.0117 | 0.6095 | 0.6780 | 21.3751 | 0.7632 | 0.7281 |

Table A.10: Evaluation results for ViT-B/16 trained on the custom balanced dataset, showing Loss, Acc1, and F1 scores for each dataset split.

Text.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax loss          | 0.2254   | 0.4367      | 0.5071 | 0.1775 | 0.0263 |
| Focal loss            | 0.2210   | 0.4206      | 0.4834 | 0.1953 | 0.0263 |
| Weighted Softmax loss | 0.1284   | 0.1760      | 0.1919 | 0.1302 | 0.0263 |
| Class-balanced loss   | 0.0558   | 0.0076      | 0.0000 | 0.0237 | 0.1053 |
| Balanced Softmax loss | 0.2460   | 0.4244      | 0.4822 | 0.2130 | 0.0789 |
| Equalization loss     | 0.2168   | 0.4215      | 0.4893 | 0.1716 | 0.0263 |
| LDAM loss             | 0.5906   | 0.6013      | 0.5924 | 0.6095 | 0.7632 |

Table A.11: Evaluation results for ViT-B/16 trained on the long-tailed dataset, showing Acc1.

| Loss Function    | Balanced |        |        | Long-tailed |        |        | Head    |        |        | Middle  |        |        | Tail    |        |        |
|------------------|----------|--------|--------|-------------|--------|--------|---------|--------|--------|---------|--------|--------|---------|--------|--------|
|                  | Loss     | Acc1   | F1     | Loss        | Acc1   | F1     | Loss    | Acc1   | F1     | Loss    | Acc1   | F1     | Loss    | Acc1   | F1     |
| Softmax          | 13.5272  | 0.2254 | 0.1871 | 6.7999      | 0.4367 | 0.4216 | 5.3024  | 0.5071 | 0.5248 | 11.3663 | 0.1775 | 0.2303 | 19.7511 | 0.0263 | 0.0263 |
| Focal Loss       | 7.5701   | 0.2210 | 0.1850 | 3.6474      | 0.4206 | 0.4016 | 2.8064  | 0.4834 | 0.4914 | 6.1246  | 0.1953 | 0.2666 | 11.3091 | 0.0263 | 0.0263 |
| Weighted Softmax | 6.5391   | 0.1284 | 0.1144 | 3.9782      | 0.1760 | 0.1902 | 3.4559  | 0.1919 | 0.2357 | 4.7288  | 0.1302 | 0.1541 | 11.0975 | 0.0263 | 0.0351 |
| Class-balanced   | 4.9938   | 0.0558 | 0.0368 | 5.8487      | 0.0076 | 0.0028 | 6.2065  | 0.0000 | 0.0000 | 4.7503  | 0.0237 | 0.0292 | 4.0694  | 0.1053 | 0.0746 |
| Balanced Softmax | 13.3583  | 0.2460 | 0.2123 | 6.7016      | 0.4244 | 0.4175 | 5.2929  | 0.4822 | 0.5121 | 11.3472 | 0.2130 | 0.2710 | 17.3287 | 0.0789 | 0.0877 |
| Equalization     | 13.4511  | 0.2168 | 0.1786 | 6.7202      | 0.4215 | 0.4062 | 5.2340  | 0.4893 | 0.5051 | 11.6755 | 0.1716 | 0.2353 | 17.4650 | 0.0263 | 0.0263 |
| LDAM             | 48.2745  | 0.5906 | 0.5926 | 47.4149     | 0.6013 | 0.6348 | 49.6692 | 0.5924 | 0.6790 | 42.0117 | 0.6095 | 0.6780 | 21.3751 | 0.7632 | 0.7281 |

Table A.12: Evaluation results for ViT-B/16 trained on the long-tailed dataset, showing Loss, Acc1, and F1 scores for each dataset split.

## A.4 ConvNeXt Base

*ConvNeXt Base trained on custom balanced dataset and imbalanced dataset on different loss functions.*

Table A.13 show the top 1 accuracies for ConvNeXt Base on various loss functions. Table A.14 show the loss, top 1 accuracy, and F1 score.

Table A.15 show the top 1 accuracies for ConvNeXt Base on various loss functions. Table A.16 show the loss, top 1 accuracy, and F1 score.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax loss          | 0.8332   | 0.8535      | 0.8566 | 0.8166 | 0.9474 |
| Focal loss            | 0.8314   | 0.8487      | 0.8507 | 0.8284 | 0.8947 |
| Weighted Softmax loss | 0.8332   | 0.8535      | 0.8566 | 0.8166 | 0.9474 |
| Class-balanced loss   | 0.8332   | 0.8535      | 0.8566 | 0.8166 | 0.9474 |
| Balanced Softmax loss | 0.8364   | 0.8344      | 0.8365 | 0.7988 | 0.9474 |
| Equalization loss     | 0.8318   | 0.8468      | 0.8448 | 0.8343 | 0.9474 |
| LDAM loss             | data     | data        | data   | data   | data   |

Table A.13: Evaluation results for ConvNeXt Base trained on the custom balanced dataset, showing Acc1.

| Loss Function    | Balanced |        |        | Long-tailed |        |        | Head   |        |        | Middle |        |        | Tail   |        |        |
|------------------|----------|--------|--------|-------------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|
|                  | Loss     | Acc1   | F1     | Loss        | Acc1   | F1     | Loss   | Acc1   | F1     | Loss   | Acc1   | F1     | Loss   | Acc1   | F1     |
| Softmax          | 0.9904   | 0.8332 | 0.8323 | 0.9594      | 0.8535 | 0.8661 | 0.9571 | 0.8566 | 0.9010 | 1.1028 | 0.8166 | 0.8603 | 0.3731 | 0.9474 | 0.9386 |
| Focal Loss       | 0.5686   | 0.8314 | 0.8301 | 0.5597      | 0.8487 | 0.8608 | 0.5640 | 0.8507 | 0.8975 | 0.6046 | 0.8284 | 0.8730 | 0.2629 | 0.8947 | 0.8947 |
| Weighted Softmax | 0.9904   | 0.8332 | 0.8323 | 0.9594      | 0.8535 | 0.8661 | 0.9571 | 0.8566 | 0.9010 | 1.1028 | 0.8166 | 0.8603 | 0.3731 | 0.9474 | 0.9386 |
| Class-balanced   | 0.9904   | 0.8332 | 0.8323 | 0.9594      | 0.8535 | 0.8661 | 0.9571 | 0.8566 | 0.9010 | 1.1028 | 0.8166 | 0.8603 | 0.3731 | 0.9474 | 0.9386 |
| Balanced Softmax | 1.0008   | 0.8364 | 0.8350 | 0.9829      | 0.8344 | 0.8478 | 0.9720 | 0.8365 | 0.8853 | 1.1509 | 0.7988 | 0.8418 | 0.4780 | 0.9474 | 0.9386 |
| Equalization     | 0.9124   | 0.8318 | 0.8302 | 0.9030      | 0.8468 | 0.8594 | 0.8550 | 0.8448 | 0.8899 | 1.2187 | 0.8343 | 0.8779 | 0.4981 | 0.9474 | 0.9474 |
| LDAM             | data     | data   | data   | data        | data   | data   | data   | data   | data   | data   | data   | data   | data   | data   | data   |

Table A.14: Evaluation results for ConvNeXt Base trained on the custom balanced dataset, showing Loss, Acc1, and F1 scores for each dataset split.

Text.

| Loss Function         | Balanced | Long-tailed | Head   | Middle | Tail   |
|-----------------------|----------|-------------|--------|--------|--------|
| Softmax loss          | 0.5972   | 0.8316      | 0.8898 | 0.6568 | 0.3158 |
| Focal loss            | 0.5938   | 0.8145      | 0.8685 | 0.6568 | 0.3158 |
| Weighted Softmax loss | 0.4090   | 0.6356      | 0.6848 | 0.4911 | 0.1842 |
| Class-balanced loss   | 0.0142   | 0.0019      | 0.0000 | 0.0000 | 0.0526 |
| Balanced Softmax loss | 0.6460   | 0.8230      | 0.8685 | 0.6509 | 0.5789 |
| Equalization loss     | 0.5956   | 0.8278      | 0.8768 | 0.6923 | 0.3421 |
| LDAM loss             | data     | data        | data   | data   | data   |

Table A.15: Evaluation results for ConvNeXt Base trained on the long-tailed dataset, showing Acc1.

| Loss Function    | Balanced |        |        | Long-tailed |        |        | Head   |        |        | Middle |        |        | Tail   |        |        |
|------------------|----------|--------|--------|-------------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|
|                  | Loss     | Acc1   | F1     | Loss        | Acc1   | F1     | Loss   | Acc1   | F1     | Loss   | Acc1   | F1     | Loss   | Acc1   | F1     |
| Softmax          | 2.7006   | 0.5972 | 0.5645 | 0.9552      | 0.8316 | 0.8202 | 0.5867 | 0.8898 | 0.9013 | 2.1181 | 0.6568 | 0.7177 | 3.9670 | 0.3158 | 0.3158 |
| Focal Loss       | 1.8210   | 0.5938 | 0.5615 | 0.6024      | 0.8145 | 0.8002 | 0.3485 | 0.8685 | 0.8791 | 1.3247 | 0.6568 | 0.7197 | 3.0291 | 0.3158 | 0.3070 |
| Weighted Softmax | 4.5284   | 0.4090 | 0.3763 | 2.0092      | 0.6356 | 0.6266 | 1.5444 | 0.6848 | 0.7054 | 3.1309 | 0.4911 | 0.5827 | 6.0533 | 0.1842 | 0.1930 |
| Class-balanced   | 5.0105   | 0.0142 | 0.0016 | 6.1643      | 0.0019 | 0.0000 | 6.5523 | 0.0000 | 0.0000 | 5.0450 | 0.0000 | 0.0000 | 3.4200 | 0.0526 | 0.0164 |
| Balanced Softmax | 2.6574   | 0.6460 | 0.6273 | 0.9120      | 0.8230 | 0.8237 | 0.5457 | 0.8685 | 0.8952 | 2.1945 | 0.6509 | 0.7250 | 3.3453 | 0.5789 | 0.5965 |
| Equalization     | 2.5527   | 0.5956 | 0.5586 | 0.9349      | 0.8278 | 0.8139 | 0.6192 | 0.8768 | 0.8907 | 1.9792 | 0.6923 | 0.7375 | 3.2293 | 0.3421 | 0.3404 |
| LDAM             | data     | data   | data   | data        | data   | data   | data   | data   | data   | data   | data   | data   | data   | data   | data   |

Table A.16: Evaluation results for ConvNeXt Base trained on the long-tailed dataset, showing Loss, Acc1, and F1 scores for each dataset split.